

## Random Forest for Carrot Price Prediction

### 4th Year Academic Research - Model Comparison

**Dataset:** Dambulla Market Carrot Prices (2020-2025)

**Goal:** Compare Random Forest with LSTM models

**Benchmark to beat:** Bidirectional LSTM - 19.30% MAPE

### Contents

1. Data Loading & Preprocessing
2. Feature Engineering
3. Feature Selection (Basic + Advanced with Heatmaps)
4. Data Preparation for Modeling
5. Random Forest - Baseline Model
6. Hyperparameter Tuning
7. Gradient Boosting Comparison
8. Model Comparison (vs LSTM)
9. Feature Importance Analysis
10. Visualizations & Results
11. Final Summary & Conclusions

```
# Install packages
!pip install scikit-learn matplotlib seaborn pandas numpy joblib -q
print("✅ Packages installed")
```

✅ Packages installed

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor
from sklearn.model_selection import RandomizedSearchCV, cross_val_score
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
import joblib
import warnings
warnings.filterwarnings('ignore')

plt.style.use('seaborn-v0_8')
sns.set_palette('husl')

print("✅ All packages loaded successfully")
```

✅ All packages loaded successfully

```
from google.colab import drive
drive.mount('/content/drive')
print("Successfully mounted!")
```

Mounted at /content/drive  
Successfully mounted!

### PART A: DATA LOADING & PREPROCESSING

```
csv_file_path = "/content/drive/MyDrive/RESEARCH-ALL-in-one/ALL-Data-in-one-CSV/best-dataset/dambulla_market_dataset.csv"

df = pd.read_csv(csv_file_path, parse_dates=["date"])
df = df.sort_values("date").reset_index(drop=True)
df.set_index("date", inplace=True)

print("*"*60)
print("📊 DATA LOADED")
```

```
print("*"*60)
print(f"Shape: {df.shape}")
print(f"Date range: {df.index.min()} to {df.index.max()}")
print(f"Columns: {len(df.columns)}")
print(f"Target variable: carrot_price")
print(f" Min: {df['carrot_price'].min():.2f} Rs")
print(f" Max: {df['carrot_price'].max():.2f} Rs")
print(f" Mean: {df['carrot_price'].mean():.2f} Rs")
print(f" Std: {df['carrot_price'].std():.2f} Rs")
```

```
=====
📊 DATA LOADED
=====
Shape: (2017, 46)
Date range: 2020-01-01 00:00:00 to 2025-07-11 00:00:00
Columns: 46
Target variable: carrot_price
Min: 53.00 Rs
Max: 1950.00 Rs
Mean: 236.80 Rs
Std: 177.85 Rs
```

```
# Data quality check
print("*"*60)
print("🔍 DATA QUALITY CHECK")
print("*"*60)

missing = df.isnull().sum()
if missing.sum() > 0:
    print(f"⚠️ Missing values found:")
    print(missing[missing > 0])
else:
    print("✅ No missing values")

# Check outliers
Q1 = df['carrot_price'].quantile(0.25)
Q3 = df['carrot_price'].quantile(0.75)
IQR = Q3 - Q1
outliers = df[(df['carrot_price'] < Q1 - 1.5*IQR) | (df['carrot_price'] > Q3 + 1.5*IQR)]
print(f"\n📉 Outliers: {len(outliers)} ({len(outliers)/len(df)*100:.2f}%)")

print(f"\n✅ Data quality check complete")
```

```
=====
🔍 DATA QUALITY CHECK
=====
✅ No missing values
📉 Outliers: 103 (5.11%)
✅ Data quality check complete
```

```
# Transform supply factors: 1→2 (high), -1→1 (normal), 0→0 (low)
supply_cols = [col for col in df.columns if 'supply_factor' in col]
df_transformed = df.copy()
supply_mapping = {1: 2, -1: 1, 0: 0}

for col in supply_cols:
    df_transformed[col] = df_transformed[col].map(supply_mapping)

print(f"✅ Transformed {len(supply_cols)} supply factor columns")
print(f"   Encoding: 2=HIGH, 1=NORMAL, 0=LOW")
```

**✓** Transformed 15 supply factor columns  
Encoding: 2=HIGH, 1=NORMAL, 0=LOW

## ▼ PART B: FEATURE ENGINEERING

**Note:** Random Forest doesn't use sequences like LSTM. Instead, we create lag features and rolling statistics as regular columns.

```
print("*"*60)
print("🔍 FEATURE ENGINEERING FOR RANDOM FOREST")
print("*"*60)

df_features = df_transformed.copy()
```

```

# =====
# 1. PRICE LAG FEATURES
# =====

print("\n1 Creating price lag features...")
df_features['price_lag_1'] = df_features['carrot_price'].shift(1)
df_features['price_lag_2'] = df_features['carrot_price'].shift(2)
df_features['price_lag_3'] = df_features['carrot_price'].shift(3)
df_features['price_lag_7'] = df_features['carrot_price'].shift(7)
df_features['price_lag_14'] = df_features['carrot_price'].shift(14)
df_features['price_lag_21'] = df_features['carrot_price'].shift(21)
df_features['price_lag_30'] = df_features['carrot_price'].shift(30)

# Rolling statistics
df_features['price_rolling_mean_7'] = df_features['carrot_price'].rolling(window=7, min_periods=1).mean()
df_features['price_rolling_mean_14'] = df_features['carrot_price'].rolling(window=14, min_periods=1).mean()
df_features['price_rolling_mean_30'] = df_features['carrot_price'].rolling(window=30, min_periods=1).mean()
df_features['price_rolling_std_7'] = df_features['carrot_price'].rolling(window=7, min_periods=1).std()
df_features['price_rolling_std_14'] = df_features['carrot_price'].rolling(window=14, min_periods=1).std()
df_features['price_rolling_min_7'] = df_features['carrot_price'].rolling(window=7, min_periods=1).min()
df_features['price_rolling_max_7'] = df_features['carrot_price'].rolling(window=7, min_periods=1).max()
df_features['price_rolling_median_7'] = df_features['carrot_price'].rolling(window=7, min_periods=1).median()

# Price changes
df_features['price_change'] = df_features['carrot_price'].diff()
df_features['price_change_pct'] = df_features['carrot_price'].pct_change()
df_features['price_change_7d'] = df_features['carrot_price'].diff(7)
df_features['price_change_14d'] = df_features['carrot_price'].diff(14)

# Volatility
df_features['price_volatility_7'] = df_features['carrot_price'].rolling(window=7).std() / df_features['carrot_price'].rolling(window=7).mean()

print(f"  Created {len([c for c in df_features.columns if 'price' in c]) - 1} price features")

# =====
# 2. PRECIPITATION FEATURES
# =====

print("\n2 Creating precipitation features...")
precip_cols = [col for col in df.columns if 'precipitation' in col]

for col in precip_cols:
    # Lag features
    df_features[f'{col}_lag_1'] = df_features[col].shift(1)
    df_features[f'{col}_lag_3'] = df_features[col].shift(3)
    df_features[f'{col}_lag_7'] = df_features[col].shift(7)
    # Rolling sum (total rain)
    df_features[f'{col}_rolling_sum_7'] = df_features[col].rolling(window=7, min_periods=1).sum()
    df_features[f'{col}_rolling_sum_14'] = df_features[col].rolling(window=14, min_periods=1).sum()

# Regional precipitation groups
base_precip_cols = [col for col in df_features.columns if 'precipitation' in col and 'lag' not in col and 'rolling' not in col]

PRECIP_GROUPS = {
    'central_highland': [col for col in base_precip_cols if any(x in col for x in ['nuwaraeliya', 'kandapola', 'ragala', 'kandy'])],
    'uva_province': [col for col in base_precip_cols if any(x in col for x in ['bandarawela', 'walmada'])],
    'northern': [col for col in base_precip_cols if 'jaffna' in col],
    'other': [col for col in base_precip_cols if not any(x in col for x in ['nuwaraeliya', 'kandapola', 'ragala', 'kandy'])]
}

for group_name, cols in PRECIP_GROUPS.items():
    if len(cols) > 0:
        df_features[f'precip_{group_name}_mean'] = df_features[cols].mean(axis=1)
        df_features[f'precip_{group_name}_max'] = df_features[cols].max(axis=1)
        df_features[f'precip_{group_name}_sum'] = df_features[cols].sum(axis=1)
        df_features[f'precip_{group_name}_mean_lag_1'] = df_features[f'precip_{group_name}_mean'].shift(1)
        df_features[f'precip_{group_name}_rolling_sum_7'] = df_features[f'precip_{group_name}_mean'].rolling(7).sum()

print(f"  Created precipitation features for {len(precip_cols)} regions and {len(PRECIP_GROUPS)} groups")

# =====
# 3. SUPPLY FACTOR FEATURES
# =====

print("\n3 Creating supply factor features...")
for col in supply_cols:
    df_features[f'{col}_lag_1'] = df_features[col].shift(1)
    df_features[f'{col}_lag_7'] = df_features[col].shift(7)
    df_features[f'{col}_rolling_mean_7'] = df_features[col].rolling(window=7, min_periods=1).mean()

```

```

df_features[f'{col}_rolling_mean_14'] = df_features[col].rolling(window=14, min_periods=1).mean()

print(f"  Created supply features for {len(supply_cols)} regions")

# =====
# 4. FUEL PRICE FEATURES
# =====
print("\n4 Creating fuel price features...")
fuel_cols = [col for col in df.columns if 'fur_' in col or any(x in col for x in ['Lp_', 'lad', 'lsd', 'lk', 'lik'])]

for col in fuel_cols:
    df_features[f'{col}_lag_1'] = df_features[col].shift(1)
    df_features[f'{col}_lag_7'] = df_features[col].shift(7)
    df_features[f'{col}_rolling_mean_7'] = df_features[col].rolling(window=7, min_periods=1).mean()

print(f"  Created fuel price features for {len(fuel_cols)} types")

# =====
# 5. TEMPORAL FEATURES
# =====
print("\n5 Creating temporal features...")
df_features['day_of_week'] = df_features.index.dayofweek
df_features['day_of_month'] = df_features.index.day
df_features['month'] = df_features.index.month
df_features['quarter'] = df_features.index.quarter
df_features['week_of_year'] = df_features.index.isocalendar().week
df_features['is_weekend'] = (df_features.index.dayofweek >= 5).astype(int)
df_features['is_month_start'] = df_features.index.is_month_start.astype(int)
df_features['is_month_end'] = df_features.index.is_month_end.astype(int)

# Cyclical encoding for temporal features
df_features['day_of_week_sin'] = np.sin(2 * np.pi * df_features['day_of_week'] / 7)
df_features['day_of_week_cos'] = np.cos(2 * np.pi * df_features['day_of_week'] / 7)
df_features['month_sin'] = np.sin(2 * np.pi * df_features['month'] / 12)
df_features['month_cos'] = np.cos(2 * np.pi * df_features['month'] / 12)

print(f"  Created temporal features")

# =====
# 6. INTERACTION FEATURES
# =====
print("\n6 Creating interaction features...")
df_features['demand_x_trading'] = df_features['dambulla_demand'] * df_features['dambulla_is_trading_activities_high_or_low']
df_features['demand_x_market_open'] = df_features['dambulla_demand'] * df_features['is_market_open']
df_features['market_open_x_weekend'] = df_features['is_market_open'] * df_features['is_weekend']

print(f"  Created interaction features")

# Fill NaN values
df_features = df_features.fillna(method='ffill').fillna(method='bfill')

# Fill any remaining NaN with median
for col in df_features.columns:
    if df_features[col].isnull().any():
        df_features[col].fillna(df_features[col].median(), inplace=True)

print("\n" + "="*60)
print("✓ FEATURE ENGINEERING COMPLETE")
print("="*60)
print(f"Total features created: {df_features.shape[1]}")
print(f"Original features: {df.shape[1]}")
print(f"New features: {df_features.shape[1] - df.shape[1]}")
print(f"Missing values: {df_features.isnull().sum().sum()}")

```

```
=====
🔧 FEATURE ENGINEERING FOR RANDOM FOREST
=====
```

- 1** Creating price lag features...
   
Created 20 price features
- 2** Creating precipitation features...
   
Created precipitation features for 17 regions and 4 groups
- 3** Creating supply factor features...
   
Created supply features for 15 regions
- 4** Creating fuel price features...

```
Created fuel price features for 9 types
```

**5** Creating temporal features...

```
Created temporal features
```

**6** Creating interaction features...

```
Created interaction features
```

```
=====
✓ FEATURE ENGINEERING COMPLETE
=====
```

```
Total features created: 273
```

```
Original features: 46
```

```
New features: 227
```

```
Missing values: 0
```

## ▼ PART C: FEATURE SELECTION (BASIC APPROACH)

**Note:** This is the basic feature selection. Advanced selection with heatmaps is in Part C.1 below.

```
print("=*60")
print("FEATURE SELECTION")
print("=*60")

# Remove rows with NaN (from lag features at the start)
# Keep last 30 days as test set before removing NaN
df_clean = df_features.copy()

# Calculate correlation with target
correlations = df_clean.corr()['carrot_price'].abs().sort_values(ascending=False)

print("\n📊 TOP 30 FEATURES BY CORRELATION:")
print(correlations.head(30))

# Quick Random Forest for feature importance
print("\n🌲 Training Random Forest for feature importance...")
X_temp = df_clean.drop('carrot_price', axis=1)
y_temp = df_clean['carrot_price']

# Remove infinite values
X_temp = X_temp.replace([np.inf, -np.inf], np.nan)
X_temp = X_temp.fillna(X_temp.median())

rf_temp = RandomForestRegressor(n_estimators=100, max_depth=15, random_state=42, n_jobs=-1)
rf_temp.fit(X_temp, y_temp)

feature_importance = pd.DataFrame({
    'feature': X_temp.columns,
    'rf_importance': rf_temp.feature_importances_
}).sort_values('rf_importance', ascending=False)

print("\n🏆 TOP 30 FEATURES BY RF IMPORTANCE:")
print(feature_importance.head(30))

# Combine approaches
corr_df = pd.DataFrame({'feature': correlations.index, 'correlation': correlations.values})
feature_scores = feature_importance.merge(corr_df, on='feature')
feature_scores['combined_score'] = (feature_scores['rf_importance'] * 0.6) + (feature_scores['correlation'] * 0.4)
feature_scores = feature_scores.sort_values('combined_score', ascending=False)

# Select features
# Strategy: Use more features for RF (it handles them well)
n_features_to_select = 50 # RF can handle more features than LSTM

selected_features = feature_scores['feature'].head(n_features_to_select).tolist()
selected_features = [f for f in selected_features if f != 'carrot_price']

# Remove highly correlated features (multicollinearity)
X_candidates = df_clean[selected_features]
candidates_corr = X_candidates.corr().abs()

features_to_remove = set()
for i in range(len(candidates_corr.columns)):
    for j in range(i+1, len(candidates_corr.columns)):
        if candidates_corr.iloc[i, j] > 0.95: # Very high correlation
            col_i = candidates_corr.columns[i]
```

```

col_j = candidates_corr.columns[]]
# Keep the one with higher target correlation
corr_i = abs(df_clean[col_i].corr(df_clean['carrot_price']))
corr_j = abs(df_clean[col_j].corr(df_clean['carrot_price']))
if corr_i < corr_j:
    features_to_remove.add(col_i)
else:
    features_to_remove.add(col_j)

final_features = [f for f in selected_features if f not in features_to_remove]

print(f"\n❌ Removed {len(features_to_remove)} highly correlated features")
print(f"✅ FINAL FEATURES: {len(final_features)}")

# Category breakdown
categories = {
    'Price Features': [f for f in final_features if 'price' in f.lower()],
    'Market Features': [f for f in final_features if any(x in f.lower() for x in ['market', 'demand', 'trading', 'dambulla']),
    'Weather Features': [f for f in final_features if 'precip' in f.lower()],
    'Supply Features': [f for f in final_features if 'supply' in f.lower()],
    'Fuel Features': [f for f in final_features if any(x in f.lower() for x in ['fur', 'lp_', 'lad', 'lsd', 'lk'])],
    'Temporal Features': [f for f in final_features if any(x in f.lower() for x in ['day', 'month', 'quarter', 'weekend', 'w'])]
}

print("\n📁 FEATURE CATEGORIES:")
for cat, feats in categories.items():
    print(f"  {cat}: {len(feats)} features")

# Save feature importance
feature_scores[feature_scores['feature'].isin(final_features)].to_csv('/content/rf_feature_importance.csv', index=False)
print("\n✅ Feature importance saved to 'rf_feature_importance.csv'")

```

---

### ⌚ FEATURE SELECTION

---

📊 TOP 30 FEATURES BY CORRELATION:

Feature	Correlation
carrot_price	1.000000
price_lag_1	0.960605
price_rolling_mean_7	0.952164
price_rolling_max_7	0.946407
price_rolling_min_7	0.937816
price_rolling_median_7	0.935490
price_lag_2	0.928781
price_rolling_mean_14	0.922189
price_lag_3	0.902969
price_lag_7	0.865969
price_rolling_mean_30	0.830756
price_rolling_std_14	0.821104
price_lag_14	0.773479
price_rolling_std_7	0.720441
price_lag_21	0.645538
price_lag_30	0.520868
price_change_14d	0.337815
lsd	0.324175
lsd_lag_1	0.323366
lsd_rolling_mean_7	0.322512
lsd_lag_7	0.319677
Lp_95	0.314247
Lp_95_rolling_mean_7	0.314070
Lp_95_lag_1	0.314019
Lp_95_lag_7	0.312993
Lp_92	0.307039
Lp_92_lag_1	0.306769
Lp_92_rolling_mean_7	0.306562
lad	0.306155
lad_lag_1	0.305529

Name: carrot\_price, dtype: float64

🌲 Training Random Forest for feature importance...

🏆 TOP 30 FEATURES BY RF IMPORTANCE:

Rank	Feature	rf_importance
58	price_rolling_max_7	0.368707
45	price_lag_1	0.265318
52	price_rolling_mean_7	0.181808
57	price_rolling_min_7	0.041184
193	hanguranketha_supply_factor_rolling_mean_14	0.012553
209	mandaramnuwara_supply_factor_rolling_mean_14	0.010919
61	price_change_pct	0.009937
62	price_change_7d	0.008370
60	price_change	0.008102

## ▼ PART C.1: ADVANCED FEATURE SELECTION WITH HEATMAPS

**Goal:** Use proper techniques including:

1. Correlation heatmap visualization
  2. Mutual Information scoring
  3. Combined feature importance
  4. Multicollinearity removal
  5. Recursive Feature Elimination (RFE)
  6. SelectFromModel for stability

```
print("=*80")
print("📸 CELL 13 - SCREENSHOT 1: CORRELATION HEATMAP")
print("🔥 STEP 1: CORRELATION HEATMAP - TOP FEATURES")
print("=*80")

# Calculate correlation matrix
corr_matrix = df_clean.corr()
target_corr = corr_matrix['carrot_price'].abs().sort_values(ascending=False)

# Select top 40 features by correlation with target
top_40_features = target_corr.head(41).index.tolist() # 41 includes target itself

print(f"\n📊 Top 40 Features by Correlation with Target:")
print(target_corr.head(40))

# Create correlation heatmap for top features
plt.figure(figsize=(20, 16))
sns.heatmap(
    df_clean[top_40_features].corr(),
    cmap='RdBu_r',
    center=0,
    annot=False,
    fmt=".2f",
    vmax=1,
    vmin=-1,
    linewidths=0.5,
    cbar_kws={'label': 'Correlation Coefficient'}
)
plt.title('Correlation Matrix - Top 40 Features (by correlation with carrot_price)')
    fontsize=16, fontweight='bold', pad=20)
plt.xticks(rotation=90, ha='right')
plt.yticks(rotation=0)
plt.tight_layout()
plt.savefig('/content/correlation_heatmap_top40.png', dpi=300, bbox_inches='tight')
plt.show()

print("\n✅ Correlation heatmap saved to 'correlation_heatmap_top40.png'")
```



```
=====
CELL 13 - SCREENSHOT 1: CORRELATION HEATMAP
STEP 1: CORRELATION HEATMAP - TOP FEATURES
=====
```

Top 40 Features by Correlation with Target:

Feature	Correlation Value
carrot_price	1.000000
price_lag_1	0.960605
price_rolling_mean_7	0.952164
price_rolling_max_7	0.946407
price_rolling_min_7	0.937816
price_rolling_median_7	0.935490
price_lag_2	0.928781
price_rolling_mean_14	0.922189
price_lag_3	0.902969
price_lag_7	0.865969
price_rolling_mean_30	0.830756
price_rolling_std_14	0.821104
price_lag_14	0.773479
price_rolling_std_7	0.720441
price_lag_21	0.645538
price_lag_30	0.520868
price_change_14d	0.337815
lsd	0.324175
lsd_lag_1	0.323366
lsd_rolling_mean_7	0.322512
lsd_lag_7	0.319677
Lp_95	0.314247
Lp_95_rolling_mean_7	0.314070
Lp_95_lag_1	0.314019
Lp_95_lag_7	0.312993

```
print("*"*80)
```

```
print("CELL 14 - SCREENSHOT 2: COMBINED FEATURE SCORES")
```

```
print("STEP 2: MUTUAL INFORMATION + RF IMPORTANCE (COMBINED SCORING)")
```

```
print("*"*80)
```

```
from sklearn.feature_selection import mutual_info_regression
```

```
# Prepare data
```

```
X_all = df_clean.drop(columns=['carrot_price']).replace([np.inf, -np.inf], np.nan).fillna(df_clean.median())
y_all = df_clean['carrot_price']
```

```
print(f"\nDataset shape: {X_all.shape}")
```

```
print(f"    Features: {X_all.shape[1]}")
```

```
print(f"    Samples: {X_all.shape[0]}")
```

```
# 1. MUTUAL INFORMATION
```

```
print("\n⌚ Computing Mutual Information scores...")
```

```
mi_scores = mutual_info_regression(X_all, y_all, random_state=42, n_neighbors=5)
```

```
mi_series = pd.Series(mi_scores, index=X_all.columns).sort_values(ascending=False)
```

```
print("\n🏆 TOP 20 FEATURES BY MUTUAL INFORMATION:")
```

```
print(mi_series.head(20))
```

```
# 2. RANDOM FOREST IMPORTANCE (more robust with more estimators)
```

```
print("\n⌚ Computing Random Forest importance...")
```

```
rf_selector = RandomForestRegressor(n_estimators=200, max_depth=15, random_state=42, n_jobs=-1)
```

```
rf_selector.fit(X_all, y_all)
```

```
rf_importance = pd.Series(rf_selector.feature_importances_, index=X_all.columns).sort_values(ascending=False)
```

```
print("\n🏆 TOP 20 FEATURES BY RF IMPORTANCE:")
```

```
print(rf_importance.head(20))
```

```
# 3. COMBINE SCORES (Weighted: RF 60%, MI 40%)
```

```
print("\n⌚ Combining scores...")
```

```
feature_scores_advanced = pd.DataFrame({
```

```
    'feature': X_all.columns,
```

```
    'mi_score': mi_series,
```

```
    'rf_importance': rf_importance,
```

```
    'correlation': target_corr[X_all.columns]
```

```
})
```

```
# Normalize all scores to 0-1 range
```

```
for col in ['mi_score', 'rf_importance', 'correlation']:
```

```
    feature_scores_advanced[f'{col}_norm'] = (
```

```
        feature_scores_advanced[col] - feature_scores_advanced[col].min() /
```

```
        feature_scores_advanced[col].max() - feature_scores_advanced[col].min())
```

```
,
```

```

# Combined score: RF(60%) + MI(30%) + Correlation(10%)
feature_scores_advanced['combined_score'] =
    feature_scores_advanced['rf_importance_norm'] * 0.60 +
    feature_scores_advanced['mi_score_norm'] * 0.30 +
    feature_scores_advanced['correlation_norm'] * 0.10
)

feature_scores_advanced = feature_scores_advanced.sort_values('combined_score', ascending=False)

print("\n🏆 TOP 30 FEATURES BY COMBINED SCORE:")
print(feature_scores_advanced[['feature', 'combined_score', 'rf_importance', 'mi_score', 'correlation']].head(30).to_string())

# Save to CSV
feature_scores_advanced.to_csv('/content/advanced_feature_scores.csv', index=False)
print("\n✅ Feature scores saved to 'advanced_feature_scores.csv'!")

```

⌚ Computing Random Forest importance...

🏆 TOP 20 FEATURES BY RF IMPORTANCE:

price_rolling_max_7	0.339112
price_lag_1	0.280064
price_rolling_mean_7	0.199232
price_rolling_min_7	0.033621
price_change_pct	0.010014
hanguranketha_supply_factor_rolling_mean_14	0.009280
price_rolling_median_7	0.008722
mandaramnuwara_supply_factor_rolling_mean_14	0.008421
price_change	0.008304
yatawatththa_supply_factor_rolling_mean_14	0.008275
price_change_7d	0.007730
puttalama_supply_factor_rolling_mean_14	0.006751
kandy_supply_factor_rolling_mean_14	0.006653
marassana_supply_factor_rolling_mean_14	0.006567
thalawakale_supply_factor_rolling_mean_14	0.006285
pussellawa_supply_factor_rolling_mean_14	0.006003
kandapola_supply_factor_rolling_mean_14	0.005775
price_lag_2	0.004831
mathale_supply_factor_rolling_mean_14	0.004124
price_change_14d	0.003989

dtype: float64

⌚ Combining scores...

🏆 TOP 30 FEATURES BY COMBINED SCORE:

	feature	combined_score	rf_importance	mi_score	correlation
	pussellawa_supply_factor_lag_7	0.991860	0.339112	1.777204	0.946407
	bandarawela_supply_factor_rolling_mean_7	0.895525	0.280064	1.817541	0.960605
	yatawatththa_supply_factor_lag_1	0.684641	0.199232	1.411720	0.952164
	yatawatththa_supply_factor_rolling_mean_7	0.447266	0.033621	1.757917	0.937816
	yatawatththa_supply_factor_lag_7	0.371480	0.008722	1.567142	0.935490
	hanguranketha_supply_factor_lag_1	0.329596	0.004831	1.359344	0.928781
	pussellawa_supply_factor_rolling_mean_7	0.308905	0.001516	1.273693	0.922189
	hanguranketha_supply_factor_rolling_mean_7	0.275621	0.000639	1.093603	0.902969
	pussellawa_supply_factor_lag_1	0.257555	0.000582	1.008160	0.865969
	pussellawa_supply_factor_rolling_mean_14	0.251952	0.000817	0.993960	0.830756
	bandarawela_supply_factor_rolling_mean_14	0.218153	0.000197	0.832071	0.773479
	yatawatththa_supply_factor_rolling_mean_14	0.195440	0.002505	0.639600	0.821104
	hanguranketha_supply_factor_lag_7	0.182805	0.000126	0.699597	0.645538
	dambulla_demand	0.173771	0.000011	0.861561	0.304874
	is_dambulla_increase	0.173368	0.000006	0.857972	0.306769
	is_market_open	0.171284	0.000012	0.845121	0.307039
	dambulla_is_trading_activities_high_or_low	0.170151	0.000010	0.838577	0.306562
	mandaramnuwara_mean_precipitation_mm_lag_7	0.167103	0.000080	0.808224	0.324175
	mandaramnuwara_mean_precipitation_mm_rolling_sum_7	0.165542	0.000036	0.799741	0.323366
	kalpitiya_mean_precipitation_mm_lag_7	0.165410	0.000006	0.812487	0.302460
	marassana_mean_precipitation_mm_lag_1	0.164460	0.000059	0.793477	0.322512
	kalpitiya_mean_precipitation_mm_lag_1	0.163929	0.000006	0.801182	0.306155
	mandaramnuwara_mean_precipitation_mm_rolling_sum_14	0.163329	0.000027	0.788764	0.319677
	kalpitiya_mean_precipitation_mm_rolling_sum_7	0.163187	0.000017	0.797372	0.304884
	kalpitiya_mean_precipitation_mm_lag_3	0.162847	0.000006	0.795018	0.305529
	kandapola_mean_precipitation_mm	0.161726	0.000015	0.783408	0.312993
	nuwaraeliya_mean_precipitation_mm	0.157329	0.000179	0.754328	0.314070
	jaffna_mean_precipitation_mm	0.156565	0.000023	0.751411	0.314019
	bandarawela_mean_precipitation_mm	0.155514	0.000029	0.711835	0.314247

```

print("=*80)
print("📸 CELL 15 - SCREENTH 3: MULTICOLLINEARITY HEATMAP")
print("🔥 STEP 3: REMOVE HIGHLY CORRELATED FEATURES (MULTICOLLINEARITY)")
print("=*80)

```

```

# Select top 80 candidates from combined score
n_candidates = 80
candidate_features = feature_scores_advanced['feature'].head(n_candidates).tolist()

print(f"\n💡 Starting with {len(candidate_features)} candidate features")

# Compute correlation matrix among candidates
X_candidates = X_all[candidate_features]
candidates_corr = X_candidates.corr().abs()

# Visualize multicollinearity
plt.figure(figsize=(18, 15))
sns.heatmap(
    candidates_corr,
    cmap='YlOrRd',
    annot=False,
    fmt=".2f",
    vmax=1,
    vmin=0,
    linewidths=0.5,
    cbar_kws={'label': 'Absolute Correlation'}
)
plt.title(f'Feature Multicollinearity Heatmap - Top {n_candidates} Candidates',
          fontsize=16, fontweight='bold', pad=20)
plt.xticks(rotation=90, ha='right', fontsize=8)
plt.yticks(rotation=0, fontsize=8)
plt.tight_layout()
plt.savefig('/content/multicollinearity_heatmap.png', dpi=300, bbox_inches='tight')
plt.show()

# Remove highly correlated features (threshold >= 0.95)
print("⏳ Removing features with correlation >= 0.95...")
features_to_remove = set()
correlation_threshold = 0.95

for i in range(len(candidates_corr.columns)):
    for j in range(i+1, len(candidates_corr.columns)):
        if candidates_corr.iloc[i, j] >= correlation_threshold:
            col_i = candidates_corr.columns[i]
            col_j = candidates_corr.columns[j]

            # Keep the feature with higher combined score
            score_i = feature_scores_advanced[feature_scores_advanced['feature'] == col_i]['combined_score'].values[0]
            score_j = feature_scores_advanced[feature_scores_advanced['feature'] == col_j]['combined_score'].values[0]

            if score_i < score_j:
                features_to_remove.add(col_i)
            else:
                features_to_remove.add(col_j)

reduced_features = [f for f in candidate_features if f not in features_to_remove]

print(f"\n❌ Removed {len(features_to_remove)} highly correlated features")
print(f"✅ Remaining features: {len(reduced_features)}")

if len(features_to_remove) > 0:
    print(f"\n📝 Removed features:")
    for feat in sorted(features_to_remove):
        print(f"    - {feat}")

```



CELL 15 - SCREENSHOT 3: MULTICOLLINEARITY HEATMAP  
 🔥 STEP 3: REMOVE HIGHLY CORRELATED FEATURES (MULTICOLLINEARITY)

Starting with 80 candidate features

Feature Multicollinearity Heatmap - Top 80 Candidates



```
print("=*80)
print("📸 CELL 16 - SCREENSHOT 4: FINAL SELECTED FEATURES")
print("🔥 STEP 4: RECURSIVE FEATURE ELIMINATION (RFE) + SelectFromModel")
print("=*80)
```

```
from sklearn.feature_selection import RFE, SelectFromModel

X_reduced = X_all[X_all.columns[selected_features]]

# METHOD 1: SelectFromModel (keeps features above median importance)
print("\n⌚ Method 1: SelectFromModel (threshold='median')...")
sfm = SelectFromModel(
    RandomForestRegressor(n_estimators=300, max_depth=20, random_state=42, n_jobs=-1),
    threshold='median'
)
sfm.fit(X_reduced, y_all)
selected_from_model = list(X_reduced.columns[sfm.get_support()])

print(f" Selected features: {len(selected_from_model)}")
print(f" Top 10: {selected_from_model[:10]}")

# METHOD 2: RFE (Recursive Feature Elimination)
print("\n⌚ Method 2: Recursive Feature Elimination (RFE)...")
n_features_rfe = min(35, max(15, len(selected_from_model))) # Target 15-35 features
rfe_estimator = RandomForestRegressor(n_estimators=200, max_depth=15, random_state=42, n_jobs=-1)
rfe = RFE(rfe_estimator, n_features_to_select=n_features_rfe, step=1)
rfe.fit(X_reduced, y_all)
selected_rfe = list(X_reduced.columns[rfe.support_])

print(f" Selected features: {len(selected_rfe)}")
print(f" Top 10: {selected_rfe[:10]}")

# COMBINE: Take intersection for stability (features selected by BOTH methods)
selected_features_final = sorted(list(set(selected_from_model) & set(selected_rfe)))

print("\n" + "=*80)
print(f"⌚ FINAL SELECTED FEATURES (Intersection for Stability): {len(selected_features_final)}")
print("=*80)

# If intersection too small, use union as fallback
if len(selected_features_final) < 10:
    print(f"⚠️ Intersection too small ({len(selected_features_final)}), using UNION instead...")
    selected_features_final = sorted(list(set(selected_from_model) | set(selected_rfe)))
    print(f"✅ FINAL FEATURES (Union): {len(selected_features_final)}")

print("\n📋 SELECTED FEATURES:")
for i, feat in enumerate(selected_features_final, 1):
```

```

score = feature_scores_advanced[feature_scores_advanced['feature'] == feature].values[-1]
print(f" {i:2d}. {feat}:50s} (score: {score:.4f})")

# Save selected features
joblib.dump(selected_features_final, '/content/selected_features_advanced.pkl')
print("\n✅ Selected features saved to 'selected_features_advanced.pkl'")

# Update final_features variable for use in subsequent cells
final_features = selected_features_final.copy()
print(f"\n✅ Updated 'final_features' variable with {len(final_features)} advanced-selected features")

=====thalawakale_supply_factor_lag_1=====
📸 EVALUATING SCREENSHOT 5: FINAL SELECTED FEATURES
🔥 STARTING RECURSIVE FEATURE ELIMINATION (RFE) + SelectFromModel
=====

🕒 Method 1: SelectFromModel (threshold='median')...
Selected features: 24
Top 10: ['bandarawela_supply_factor_rolling_mean_7', 'pussellawa_supply_factor_rolling_mean_14', 'bandarawela_supply_fac

🕒 Method 2: Recursive Feature Elimination (RFE)...
Selected features: 24
Top 10: ['bandarawela_supply_factor_rolling_mean_7', 'pussellawa_supply_factor_rolling_mean_14', 'bandarawela_supply_fac

=====FINAL SELECTED FEATURES (Intersection for Stability): 24=====
=====

📁 SELECTED FEATURES:
1. bandarawela_supply_factor_rolling_mean_14 (score: 0.2182)
2. bandarawela_supply_factor_rolling_mean_7 (score: 0.8955)
3. fur_1500_high_rolling_mean_7 (score: 0.0468)
4. is_dambulla_increase (score: 0.1734)
5. jaffna_mean_precipitation_mm (score: 0.1566)
6. kalpitiya_mean_precipitation_mm_lag_7 (score: 0.1654)
7. kalpitiya_mean_precipitation_mm_rolling_sum_14 (score: 0.1251)
8. kalpitiya_mean_precipitation_mm_rolling_sum_7 (score: 0.1632)
9. kandapola_mean_precipitation_mm (score: 0.1617)
10. kandapola_supply_factor_rolling_mean_14 (score: 0.1421)
11. lad_lag_1 (score: 0.0608)
12. lk_lag_1 (score: 0.0460)
13. mandaramnuwara_mean_precipitation_mm_lag_3 (score: 0.1473)
14. mandaramnuwara_mean_precipitation_mm_lag_7 (score: 0.1671)
15. mandaramnuwara_mean_precipitation_mm_rolling_sum_14 (score: 0.1633)
16. mandaramnuwara_mean_precipitation_mm_rolling_sum_7 (score: 0.1655)
17. marassana_mean_precipitation_mm_lag_1 (score: 0.1645)
18. precip_central_highland_mean (score: 0.0970)
19. precip_uva_province_sum (score: 0.0446)
20. price_lag_14 (score: 0.0455)
21. price_rolling_mean_30 (score: 0.0487)
22. pussellawa_supply_factor_rolling_mean_14 (score: 0.2520)
23. quarter (score: 0.1088)
24. yatawaththa_supply_factor_rolling_mean_14 (score: 0.1954)

✅ Selected features saved to 'selected_features_advanced.pkl'
✅ Updated 'final_features' variable with 24 advanced-selected features

```

```

print("=*80")
print("📸 CELL 17 - SCREENSHOT 5: FEATURE SELECTION SUMMARY DASHBOARD")
print("📊 FEATURE SELECTION SUMMARY & VISUALIZATION")
print("=*80")

# Create comparison visualization
fig, axes = plt.subplots(2, 2, figsize=(18, 12))

# 1. Top 20 features by combined score
top_20_combined = feature_scores_advanced.head(20)
axes[0, 0].barh(range(len(top_20_combined)), top_20_combined['combined_score'], color='#2E86AB')
axes[0, 0].set_yticks(range(len(top_20_combined)))
axes[0, 0].set_yticklabels(top_20_combined['feature'], fontsize=9)
axes[0, 0].set_xlabel('Combined Score', fontsize=11)
axes[0, 0].set_title('Top 20 Features by Combined Score', fontsize=12, fontweight='bold')
axes[0, 0].invert_yaxis()
axes[0, 0].grid(alpha=0.3, axis='x')

# 2. Feature selection method comparison
methods = ['Initial\nCandidates', 'After\nMulticollinearity', 'SelectFromModel', 'RFE', 'Final\n(Intersection)']
counts = [len(candidate_features), len(reduced_features), len(selected_from_model),
         len(selected_rfe), len(selected_features_final)]
colors_bar = ['#A8DADC', '#457B00', '#1A23557', '#F63046', '#A6D6AA']

=====

```

```

    axes[0, 1].bar(range(len(methods)), counts, color=colors_bar)
    axes[0, 1].set_xticks(range(len(methods)))
    axes[0, 1].set_xticklabels(methods, fontsize=10)
    axes[0, 1].set_ylabel('Number of Features', fontsize=11)
    axes[0, 1].set_title('Feature Selection Pipeline', fontsize=12, fontweight='bold')
    axes[0, 1].grid(alpha=0.3, axis='y')
    for i, v in enumerate(counts):
        axes[0, 1].text(i, v + 1, str(v), ha='center', fontweight='bold')

# 3. Score distribution comparison
score_types = ['RF Importance', 'Mutual Info', 'Correlation', 'Combined']
score_cols = ['rf_importance', 'mi_score', 'correlation', 'combined_score']
final_features_df = feature_scores_advanced[feature_scores_advanced['feature'].isin(selected_features_final)]

bp = axes[1, 0].boxplot(
    [final_features_df[col] for col in score_cols],
    labels=score_types,
    patch_artist=True,
    showmeans=True
)
for patch, color in zip(bp['boxes'], ['#FF6B6B', '#4ECD4', '#45B7D1', '#96CEB4']):
    patch.set_facecolor(color)
axes[1, 0].set_ylabel('Score', fontsize=11)
axes[1, 0].set_title('Score Distribution of Selected Features', fontsize=12, fontweight='bold')
axes[1, 0].grid(alpha=0.3, axis='y')
axes[1, 0].tick_params(axis='x', rotation=45)

# 4. Category breakdown of selected features
categories_new = {
    'Price Features': [f for f in selected_features_final if 'price' in f.lower()],
    'Weather': [f for f in selected_features_final if 'precip' in f.lower()],
    'Supply': [f for f in selected_features_final if 'supply' in f.lower()],
    'Market': [f for f in selected_features_final if any(x in f.lower() for x in ['market', 'demand', 'trading', 'dambulla'])],
    'Fuel': [f for f in selected_features_final if any(x in f.lower() for x in ['fur', 'lp_', 'lad', 'lsd', 'lk'])],
    'Temporal': [f for f in selected_features_final if any(x in f.lower() for x in ['day', 'month', 'quarter', 'weekend', 'w'
}

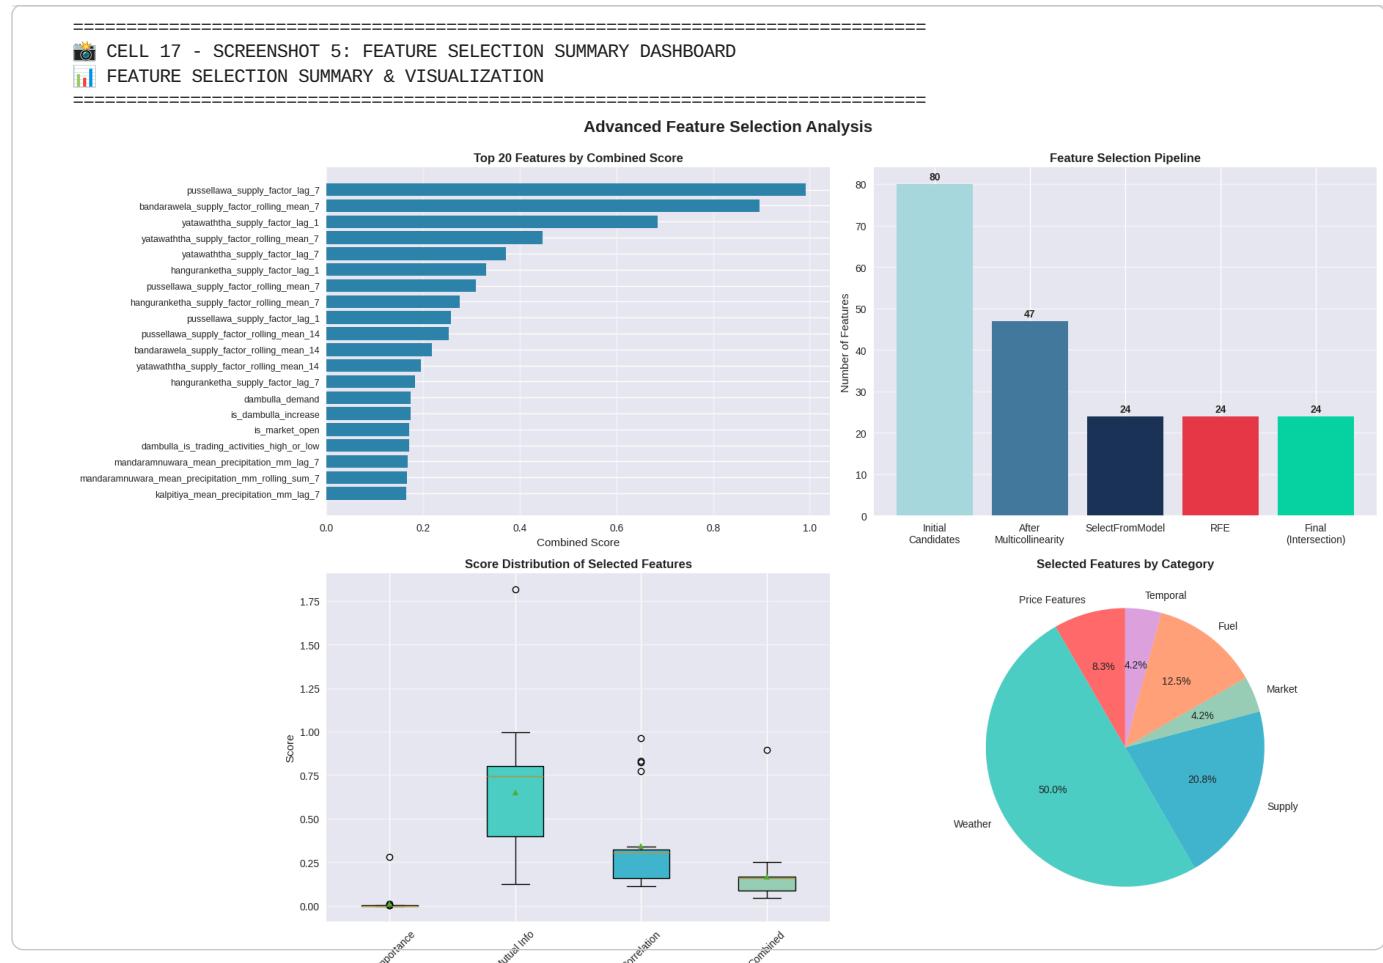
cat_names = list(categories_new.keys())
cat_counts = [len(feats) for feats in categories_new.values()]
axes[1, 1].pie(cat_counts, labels=cat_names, autopct='%1.1f%%', startangle=90,
               colors=[ '#FF6B6B', '#4ECD4', '#45B7D1', '#96CEB4', '#FFA07A', '#DDA0DD'])
axes[1, 1].set_title('Selected Features by Category', fontsize=12, fontweight='bold')

plt.suptitle('Advanced Feature Selection Analysis', fontsize=16, fontweight='bold', y=0.995)
plt.tight_layout()
plt.savefig('/content/feature_selection_summary.png', dpi=300, bbox_inches='tight')
plt.show()

print("\n📁 CATEGORY BREAKDOWN:")
for cat, feats in categories_new.items():
    print(f"  {cat}: {len(feats)} features")

print("\n✅ Feature selection summary saved to 'feature_selection_summary.png'")

```



## ▼ PART D: PREPARE DATA FOR MODELING

### CATEGORY BREAKDOWN:

Price Features: 2 features

```

print("*" * 60)
print("PREPARING DATA FOR RANDOM FOREST")
print("*" * 60)

# Create final dataset
df_final = df_clean[final_features + ['carrot_price']].copy()

# Remove any remaining NaN or inf
df_final = df_final.replace([np.inf, -np.inf], np.nan)
df_final = df_final.fillna(method='ffill').fillna(method='bfill')

for col in df_final.columns:
    if df_final[col].isnull().any():
        df_final[col].fillna(df_final[col].median(), inplace=True)

print(f"Final dataset: {df_final.shape}")
print(f"Features: {len(final_features)}")
print(f"Samples: {len(df_final)}")
print(f"Missing values: {df_final.isnull().sum().sum()}")
print(f"Infinite values: {np.isinf(df_final.select_dtypes(include=[np.number])).sum().sum()}")


# Time series split (70% train, 15% val, 15% test)
# IMPORTANT: No shuffling for time series!
train_size = int(len(df_final) * 0.70)
val_size = int(len(df_final) * 0.15)

train_data = df_final.iloc[:train_size]
val_data = df_final.iloc[train_size:train_size+val_size]
test_data = df_final.iloc[train_size+val_size:]

X_train = train_data[final_features]
y_train = train_data['carrot_price']

X_val = val_data[final_features]
y_val = val_data['carrot_price']

```

```
X_test = test_data[final_features]
y_test = test_data['carrot_price']

print("\n📊 DATA SPLIT (Time Series - No Shuffling):")
print(f" Train: {len(X_train)} samples ({len(X_train)/len(df_final)*100:.1f}%)")
print(f" Val: {len(X_val)} samples ({len(X_val)/len(df_final)*100:.1f}%)")
print(f" Test: {len(X_test)} samples ({len(X_test)/len(df_final)*100:.1f}%)")

print("\n📅 Date ranges:")
print(f" Train: {train_data.index.min()} to {train_data.index.max()}")
print(f" Val: {val_data.index.min()} to {val_data.index.max()}")
print(f" Test: {test_data.index.min()} to {test_data.index.max()}")


=====
📊 PREPARING DATA FOR RANDOM FOREST
=====

✓ Final dataset: (2017, 25)
Features: 24
Samples: 2017
Missing values: 0
Infinite values: 0

📊 DATA SPLIT (Time Series - No Shuffling):
Train: 1411 samples (70.0%)
Val: 302 samples (15.0%)
Test: 304 samples (15.1%)

📅 Date ranges:
Train: 2020-01-01 00:00:00 to 2023-11-13 00:00:00
Val: 2023-11-14 00:00:00 to 2024-09-10 00:00:00
Test: 2024-09-11 00:00:00 to 2025-07-11 00:00:00
```

## ▼ PART E: RANDOM FOREST - BASELINE MODEL

```
def calc_mape(actual, pred):
    return np.mean(np.abs((actual - pred) / actual)) * 100

def evaluate_model(model, X_train, y_train, X_val, y_val, X_test, y_test, model_name):
    # Predictions
    y_train_pred = model.predict(X_train)
    y_val_pred = model.predict(X_val)
    y_test_pred = model.predict(X_test)

    # Metrics
    results = {
        'model': model_name,
        'train_mape': calc_mape(y_train, y_train_pred),
        'val_mape': calc_mape(y_val, y_val_pred),
        'test_mape': calc_mape(y_test, y_test_pred),
        'test_mae': mean_absolute_error(y_test, y_test_pred),
        'test_rmse': np.sqrt(mean_squared_error(y_test, y_test_pred)),
        'test_r2': r2_score(y_test, y_test_pred),
        'predictions': {
            'train': (y_train, y_train_pred),
            'val': (y_val, y_val_pred),
            'test': (y_test, y_test_pred)
        }
    }

    return results

print("*"*60)
print("📸 CELL 20 - SCREENSHOT 6: BASELINE MODEL RESULTS")
print("🌲 RANDOM FOREST - BASELINE MODEL")
print("*"*60)

# Baseline Random Forest with default parameters
rf_baseline = RandomForestRegressor(
    n_estimators=100,
    random_state=42,
    n_jobs=-1,
    verbose=1
)

print("\n⌚ Training baseline model...")
```

```

rf_baseline.fit(X_train, y_train)

results_baseline = evaluate_model(rf_baseline, X_train, y_train, X_val, y_val, X_test, y_test, "RF Baseline")

print("\n✅ BASELINE RESULTS:")
print(f" Train MAPE: {results_baseline['train_mape']:.2f}%")
print(f" Val MAPE: {results_baseline['val_mape']:.2f}%")
print(f" Test MAPE: {results_baseline['test_mape']:.2f}%")
print(f" Test MAE: {results_baseline['test_mae']:.2f} Rs")
print(f" Test RMSE: {results_baseline['test_rmse']:.2f} Rs")
print(f" Test R²: {results_baseline['test_r2']:.4f}")

# Compare with LSTM benchmark
lstm_benchmark = 19.30
print(f"\n📊 COMPARISON TO LSTM:")
print(f" Best LSTM (Bidirectional): {lstm_benchmark:.2f}% MAPE")
print(f" RF Baseline: {results_baseline['test_mape']:.2f}% MAPE")
diff = lstm_benchmark - results_baseline['test_mape']
if diff > 0:
    print(f" ✅ RF is BETTER by {diff:.2f}% points!")
elif diff < 0:
    print(f" ⚠️ LSTM is better by {abs(diff):.2f}% points")
else:
    print(f" Same performance")

# Save baseline model
joblib.dump(rf_baseline, '/content/rf_baseline_model.pkl')
print("\n✅ Baseline model saved")

```

```
=====
📸 CELL 20 - SCREENSHOT 6: BASELINE MODEL RESULTS
🎄 RANDOM FOREST - BASELINE MODEL
=====
```

```
⌚ Training baseline model...
[Parallel(n_jobs=-1)]: Using backend ThreadingBackend with 2 concurrent workers.
[Parallel(n_jobs=-1)]: Done 46 tasks | elapsed: 0.7s
```

#### ✅ BASELINE RESULTS:

```

Train MAPE: 4.14%
Val MAPE: 29.53%
Test MAPE: 34.13%
Test MAE: 124.40 Rs
Test RMSE: 179.98 Rs
Test R²: 0.3800

```

#### 📊 COMPARISON TO LSTM:

```

Best LSTM (Bidirectional): 19.30% MAPE
RF Baseline: 34.13% MAPE
⚠️ LSTM is better by 14.83% points

```

#### ✅ Baseline model saved

```

[Parallel(n_jobs=-1)]: Done 100 out of 100 | elapsed: 1.5s finished
[Parallel(n_jobs=2)]: Using backend ThreadingBackend with 2 concurrent workers.
[Parallel(n_jobs=2)]: Done 46 tasks | elapsed: 0.0s
[Parallel(n_jobs=2)]: Done 100 out of 100 | elapsed: 0.0s finished
[Parallel(n_jobs=2)]: Using backend ThreadingBackend with 2 concurrent workers.
[Parallel(n_jobs=2)]: Done 46 tasks | elapsed: 0.0s
[Parallel(n_jobs=2)]: Done 100 out of 100 | elapsed: 0.0s finished
[Parallel(n_jobs=2)]: Using backend ThreadingBackend with 2 concurrent workers.
[Parallel(n_jobs=2)]: Done 46 tasks | elapsed: 0.0s
[Parallel(n_jobs=2)]: Done 100 out of 100 | elapsed: 0.0s finished

```

## ▼ PART F: HYPERPARAMETER TUNING

```

print("=*60)
print("📸 CELL 22 - SCREENSHOT 7: TUNED MODEL RESULTS")
print("🎯 HYPERPARAMETER TUNING")
print("=*60)

# Define parameter grid
param_distributions = {
    'n_estimators': [100, 200, 300, 500],
    'max_depth': [10, 15, 20, 25, 30, None],
    'min_samples_split': [2, 5, 10, 15],
    'min_samples_leaf': [1, 2, 4, 8],
    'max_features': ['sqrt', 'log2', 0.5, 0.7],
    'bootstrap': [True, False]
}

```

```

}

print("\n🔍 Parameter search space:")
for param, values in param_distributions.items():
    print(f"    {param}: {values}")

# Randomized search (faster than grid search)
rf_random = RandomForestRegressor(random_state=42, n_jobs=-1)

random_search = RandomizedSearchCV(
    estimator=rf_random,
    param_distributions=param_distributions,
    n_iter=50, # Try 50 different combinations
    cv=3, # 3-fold cross-validation
    scoring='neg_mean_absolute_error',
    random_state=42,
    n_jobs=-1,
    verbose=2
)

print("\n⌚ Starting hyperparameter search (this may take a few minutes)...")
random_search.fit(X_train, y_train)

print("\n✅ BEST PARAMETERS FOUND:")
for param, value in random_search.best_params_.items():
    print(f"    {param}: {value}")

# Best model
rf_tuned = random_search.best_estimator_

results_tuned = evaluate_model(rf_tuned, X_train, y_train, X_val, y_val, X_test, y_test, "RF Tuned")

print("\n✅ TUNED MODEL RESULTS:")
print(f"    Train MAPE: {results_tuned['train_mape']:.2f}%")
print(f"    Val MAPE: {results_tuned['val_mape']:.2f}%")
print(f"    Test MAPE: {results_tuned['test_mape']:.2f}%")
print(f"    Test MAE: {results_tuned['test_mae']:.2f} Rs")
print(f"    Test RMSE: {results_tuned['test_rmse']:.2f} Rs")
print(f"    Test R²: {results_tuned['test_r2']:.4f}")

improvement = results_baseline['test_mape'] - results_tuned['test_mape']
print("\n⌚ IMPROVEMENT:")
print(f"    Baseline: {results_baseline['test_mape']:.2f}% MAPE")
print(f"    Tuned: {results_tuned['test_mape']:.2f}% MAPE")
print(f"    Gain: {improvement:.2f} points ({(improvement/results_baseline['test_mape'])*100:.1f}% relative)")

# Save tuned model
joblib.dump(rf_tuned, '/content/rf_tuned_model.pkl')
print("\n✅ Tuned model saved")

```

=====  
📸 CELL 22 - SCREENSHOT 7: TUNED MODEL RESULTS  
⌚ HYPERPARAMETER TUNING  
=====

🔍 Parameter search space:  
n\_estimators: [100, 200, 300, 500]  
max\_depth: [10, 15, 20, 25, 30, None]  
min\_samples\_split: [2, 5, 10, 15]  
min\_samples\_leaf: [1, 2, 4, 8]  
max\_features: ['sqrt', 'log2', 0.5, 0.7]  
bootstrap: [True, False]

⌚ Starting hyperparameter search (this may take a few minutes)...  
Fitting 3 folds for each of 50 candidates, totalling 150 fits

✅ BEST PARAMETERS FOUND:  
n\_estimators: 100  
min\_samples\_split: 2  
min\_samples\_leaf: 8  
max\_features: 0.7  
max\_depth: 10  
bootstrap: True

✅ TUNED MODEL RESULTS:  
Train MAPE: 9.82%  
Val MAPE: 27.65%  
Test MAPE: 34.10%  
Test MAE: 123.43 Rs  
Test RMSE: 178.08 Rs

Test R<sup>2</sup>: 0.3931

## IMPROVEMENT:

Baseline: 34.13% MAPE  
 Tuned: 34.10% MAPE  
 Gain: 0.03% points (0.1% relative)

Tuned model saved

## ▼ PART G: GRADIENT BOOSTING (BONUS COMPARISON)

```

print("=*60)
print("🚀 GRADIENT BOOSTING (Bonus Comparison)")
print("=*60)

# Gradient Boosting often performs better than RF
gb_model = GradientBoostingRegressor(
    n_estimators=200,
    learning_rate=0.1,
    max_depth=5,
    min_samples_split=10,
    min_samples_leaf=4,
    subsample=0.8,
    random_state=42,
    verbose=1
)

print("\n⌚ Training Gradient Boosting...")
gb_model.fit(X_train, y_train)

results_gb = evaluate_model(gb_model, X_train, y_train, X_val, y_val, X_test, y_test, "Gradient Boosting")

print("\n✅ GRADIENT BOOSTING RESULTS:")
print(f" Train MAPE: {results_gb['train_mape']:.2f}%")
print(f" Val MAPE: {results_gb['val_mape']:.2f}%")
print(f" Test MAPE: {results_gb['test_mape']:.2f}%")
print(f" Test MAE: {results_gb['test_mae']:.2f} Rs")
print(f" Test RMSE: {results_gb['test_rmse']:.2f} Rs")
print(f" Test R2: {results_gb['test_r2']:.4f}")

# Save GB model
joblib.dump(gb_model, '/content/gb_model.pkl')
print("\n✅ Gradient Boosting model saved")

```

---

GRADIENT BOOSTING (Bonus Comparison)

---

⌚ Training Gradient Boosting...

Iter	Train Loss	OOB Improve	Remaining Time
1	7973.9950	1367.2202	3.63s
2	6495.8533	395.1366	2.88s
3	5577.0921	1594.6547	2.71s
4	4726.4311	915.7863	2.54s
5	4018.6518	663.3923	2.42s
6	3243.2833	-268.4299	2.38s
7	2794.1909	545.6988	2.34s
8	2523.8089	816.5646	2.28s
9	2253.2326	501.0364	2.24s
10	1894.4309	-83.1740	2.21s
20	738.8775	-43.4662	2.07s
30	415.4636	-53.1521	1.92s
40	310.3835	-56.6236	1.78s
50	252.7837	-28.8747	1.66s
60	201.1987	-10.4237	1.54s
70	182.3402	47.5333	1.42s
80	155.9369	11.4072	1.30s
90	137.6386	0.6709	1.19s
100	118.1954	-6.2239	1.08s
200	36.9928	7.4842	0.00s

GRADIENT BOOSTING RESULTS:

Train MAPE: 2.91%  
 Val MAPE: 27.32%  
 Test MAPE: 34.00%  
 Test MAE: 122.63 Rs  
 Test RMSE: 174.08 Rs  
 Test R<sup>2</sup>: 0.4201

Gradient Boosting model saved

## ▼ PART H: COMPREHENSIVE MODEL COMPARISON

```

print("=*80")
print("📸 CELL 25 - SCREENSHOT 8: FINAL MODEL COMPARISON TABLE")
print("📊 COMPREHENSIVE MODEL COMPARISON")
print("=*80")

# Create comparison dataframe
all_results = [results_baseline, results_tuned, results_gb]

comparison_df = pd.DataFrame([
    {
        'Model': r['model'],
        'Train MAPE': f'{r["train_mape"]:.2f}%',
        'Val MAPE': f'{r["val_mape"]:.2f}%',
        'Test MAPE': f'{r["test_mape"]:.2f}%',
        'Test MAE': f'{r["test_mae"]:.2f} Rs',
        'Test RMSE': f'{r["test_rmse"]:.2f} Rs',
        'Test R^2': f'{r["test_r2"]:.4f}'
    }
    for r in all_results
])

# Add LSTM results for comparison
lstm_comparison = pd.DataFrame([
    {
        'Model': 'LSTM (Bidirectional)',
        'Train MAPE': '13.09%',
        'Val MAPE': '14.88%',
        'Test MAPE': '19.30%',
        'Test MAE': '63.10 Rs',
        'Test RMSE': '91.99 Rs',
        'Test R^2': '0.8384'
    }
])

full_comparison = pd.concat([comparison_df, lstm_comparison], ignore_index=True)

print("\n" + full_comparison.to_string(index=False))

# Find best model
best_rf_idx = comparison_df['Test MAPE'].str.rstrip('%').astype(float).idxmin()
best_rf_name = comparison_df.loc[best_rf_idx, 'Model']
best_rf_mape = float(comparison_df.loc[best_rf_idx, 'Test MAPE'].rstrip('%'))

print(f"\n🏆 BEST RANDOM FOREST MODEL: {best_rf_name}")
print(f"    Test MAPE: {best_rf_mape:.2f}%")

print(f"\n📊 FINAL COMPARISON:")
print(f"    Best LSTM: 19.30% MAPE (R² = 0.8384)")
print(f"    Best Random Forest: {best_rf_mape:.2f}% MAPE (R² = {comparison_df.loc[best_rf_idx, 'Test R²']}%)")

diff = 19.30 - best_rf_mape
if diff > 0:
    print(f"\n🎉 WINNER: Random Forest is BETTER by {diff:.2f}% points!")
elif diff < 0:
    print(f"\n🎉 WINNER: LSTM is better by {abs(diff):.2f}% points")
else:
    print(f"\n🤝 TIE: Both models perform equally")

# Save comparison
full_comparison.to_csv('/content/model_comparison_all.csv', index=False)
print("\n✅ Comparison saved to 'model_comparison_all.csv'")

```

```
=====
📸 CELL 25 - SCREENSHOT 8: FINAL MODEL COMPARISON TABLE
📊 COMPREHENSIVE MODEL COMPARISON
=====
```

	Model	Train MAPE	Val MAPE	Test MAPE	Test MAE	Test RMSE	Test R <sup>2</sup>
RF Baseline		4.14%	29.53%	34.13%	124.40 Rs	179.98 Rs	0.3800
RF Tuned		9.82%	27.65%	34.10%	123.43 Rs	178.08 Rs	0.3931
Gradient Boosting		2.91%	27.32%	34.00%	122.63 Rs	174.08 Rs	0.4201
LSTM (Bidirectional)		13.09%	14.88%	19.30%	63.10 Rs	91.99 Rs	0.8384

🏆 BEST RANDOM FOREST MODEL: Gradient Boosting  
Test MAPE: 34.00%

📊 FINAL COMPARISON:  
Best LSTM: 19.30% MAPE ( $R^2 = 0.8384$ )  
Best Random Forest: 34.00% MAPE ( $R^2 = 0.4201$ )

🏆 WINNER: LSTM is better by 14.70% points

✓ Comparison saved to 'model\_comparison\_all.csv'

## ▼ PART I: FEATURE IMPORTANCE ANALYSIS

```
print("=*60)
print("📊 FEATURE IMPORTANCE ANALYSIS")
print("=*60)

# Get feature importance from best model
if best_rf_name == "RF Baseline":
    best_model = rf_baseline
elif best_rf_name == "RF Tuned":
    best_model = rf_tuned
else:
    best_model = gb_model

feature_imp_df = pd.DataFrame({
    'feature': final_features,
    'importance': best_model.feature_importances_
}).sort_values('importance', ascending=False)

print("\n🏆 TOP 30 MOST IMPORTANT FEATURES:")
print(feature_imp_df.head(30).to_string(index=False))

# Category-wise importance
category_importance = {}
for cat, feats in categories.items():
    cat_features = [f for f in feats if f in final_features]
    if cat_features:
        cat_imp = feature_imp_df[feature_imp_df['feature'].isin(cat_features)]['importance'].sum()
        category_importance[cat] = cat_imp

print("\n📊 IMPORTANCE BY CATEGORY:")
for cat, imp in sorted(category_importance.items(), key=lambda x: x[1], reverse=True):
    print(f"  {cat}: {imp:.4f} ({imp/sum(category_importance.values())*100:.1f}%)")

# Save feature importance
feature_imp_df.to_csv('/content/best_model_feature_importance.csv', index=False)
print("\n✓ Feature importance saved")
```

=====

📊 FEATURE IMPORTANCE ANALYSIS

=====

🏆 TOP 30 MOST IMPORTANT FEATURES:

	feature	importance
	price_rolling_mean_30	0.753500
	quarter	0.025708
kalpitiya_mean_precipitation_mm_rolling_sum_14		0.025070
mandaramnuwara_mean_precipitation_mm_rolling_sum_14		0.024988
kandapola_supply_factor_rolling_mean_14		0.020801
is_dambulla_increase		0.016121
fur_1500_high_rolling_mean_7		0.013328
bandarawela_supply_factor_rolling_mean_14		0.011795
pussellawa_supply_factor_rolling_mean_14		0.011544
kalpitiya_mean_precipitation_mm_rolling_sum_7		0.009780
mandaramnuwara_mean_precipitation_mm_rolling_sum_7		0.009662
price_lag_14		0.008084
lad_lag_1		0.007990
lk_lag_1		0.007360
bandarawela_supply_factor_rolling_mean_7		0.006782
mandaramnuwara_mean_precipitation_mm_lag_3		0.006740
kandapola_mean_precipitation_mm		0.006257
kalpitiya_mean_precipitation_mm_lag_7		0.006032
mandaramnuwara_mean_precipitation_mm_lag_7		0.005601
precip_uva_province_sum		0.005523
yatawatththa_supply_factor_rolling_mean_14		0.004976
precip_central_highland_mean		0.004526
jaffna_mean_precipitation_mm		0.003946

marassana\_mean\_precipitation\_mm\_lag\_1 0.003887

IMPORTANCE BY CATEGORY:  
 Temporal Features: 0.0257 (76.1%)  
 Price Features: 0.0081 (23.9%)

Feature importance saved

## ▼ PART J: VISUALIZATIONS

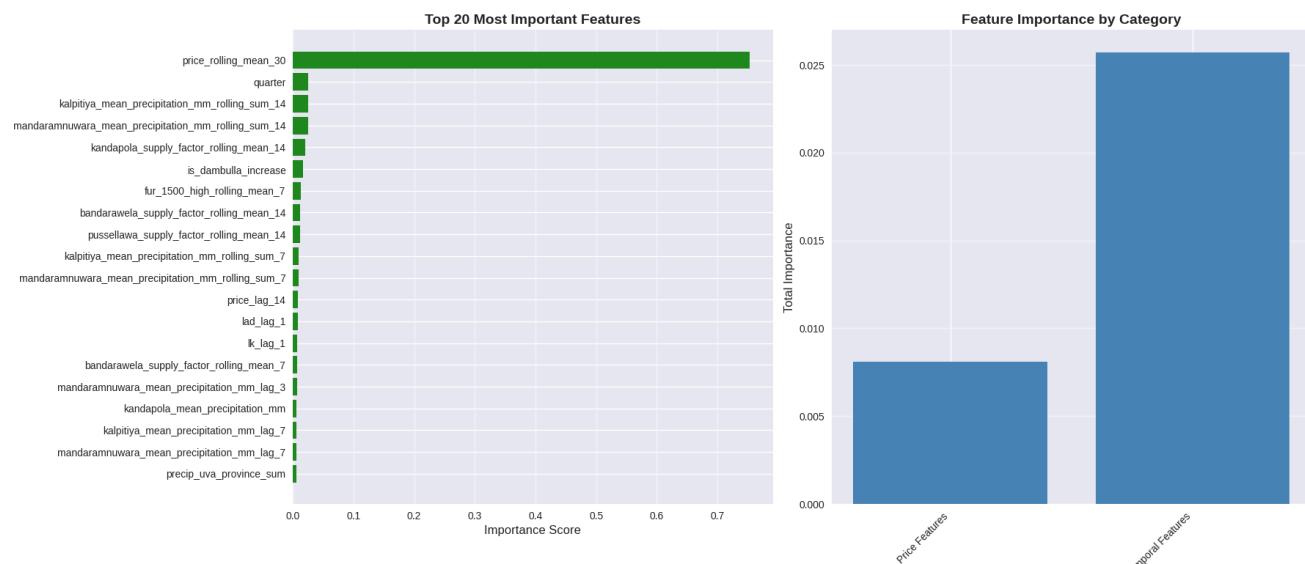
```
# 1. Feature Importance Plot
fig, axes = plt.subplots(1, 2, figsize=(18, 8))

# Top 20 features
top_20 = feature_imp_df.head(20)
axes[0].barh(range(len(top_20)), top_20['importance'], color='forestgreen')
axes[0].set_yticks(range(len(top_20)))
axes[0].set_yticklabels(top_20['feature'])
axes[0].set_xlabel('Importance Score', fontsize=12)
axes[0].set_title('Top 20 Most Important Features', fontsize=14, fontweight='bold')
axes[0].invert_yaxis()
axes[0].grid(alpha=0.3, axis='x')

# Category importance
cat_names = list(category_importance.keys())
cat_values = list(category_importance.values())
axes[1].bar(range(len(cat_names)), cat_values, color='steelblue')
axes[1].set_xticks(range(len(cat_names)))
axes[1].set_xticklabels(cat_names, rotation=45, ha='right')
axes[1].set_ylabel('Total Importance', fontsize=12)
axes[1].set_title('Feature Importance by Category', fontsize=14, fontweight='bold')
axes[1].grid(alpha=0.3, axis='y')

plt.tight_layout()
plt.savefig('/content/feature_importance.png', dpi=300, bbox_inches='tight')
plt.show()

print("✅ Feature importance visualization saved")
```



Feature importance visualization saved

```
# 2. Predictions Comparison
fig, axes = plt.subplots(3, 1, figsize=(18, 14))
```

```
# Use best RF model predictions
best_rf_results = results_tuned if best_rf_name == "RF Tuned" else results_baseline

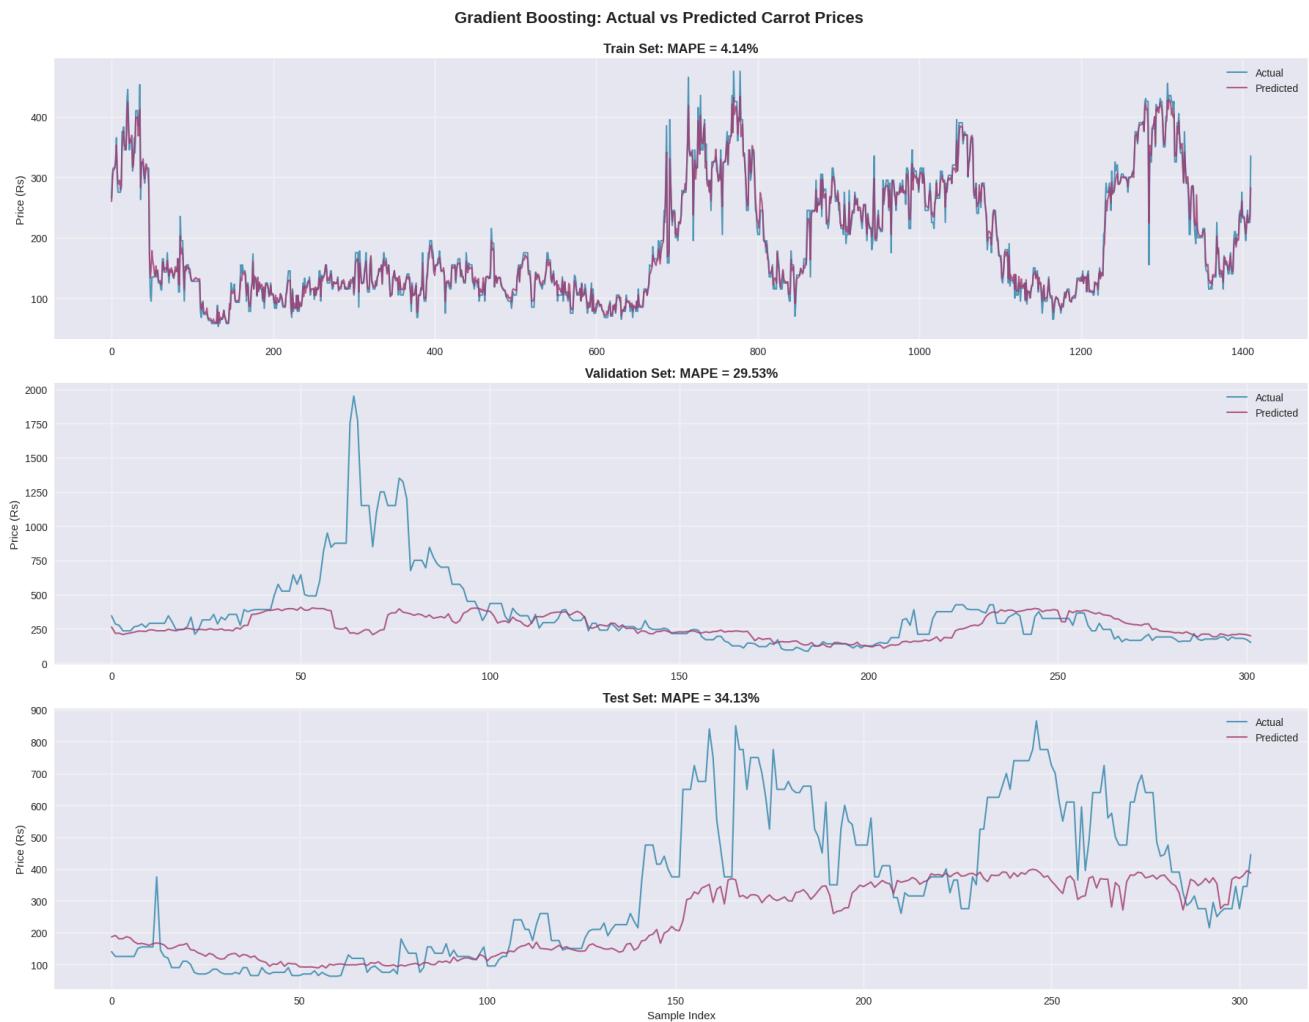
# Train
y_train_actual, y_train_pred = best_rf_results['predictions']['train']
axes[0].plot(y_train_actual.values, label='Actual', linewidth=1.5, alpha=0.8, color='#2E86AB')
axes[0].plot(y_train_pred, label='Predicted', linewidth=1.5, alpha=0.8, color="#A23B72")
axes[0].set_title(f'Train Set: MAPE = {best_rf_results["train_mape"]:.2f}%', fontsize=13, fontweight='bold')
axes[0].set_ylabel('Price (Rs)', fontsize=11)
axes[0].legend(fontsize=10)
axes[0].grid(alpha=0.3)

# Validation
y_val_actual, y_val_pred = best_rf_results['predictions']['val']
axes[1].plot(y_val_actual.values, label='Actual', linewidth=1.5, alpha=0.8, color='#2E86AB')
axes[1].plot(y_val_pred, label='Predicted', linewidth=1.5, alpha=0.8, color="#A23B72")
axes[1].set_title(f'Validation Set: MAPE = {best_rf_results["val_mape"]:.2f}%', fontsize=13, fontweight='bold')
axes[1].set_ylabel('Price (Rs)', fontsize=11)
axes[1].legend(fontsize=10)
axes[1].grid(alpha=0.3)

# Test
y_test_actual, y_test_pred = best_rf_results['predictions']['test']
axes[2].plot(y_test_actual.values, label='Actual', linewidth=1.5, alpha=0.8, color='#2E86AB')
axes[2].plot(y_test_pred, label='Predicted', linewidth=1.5, alpha=0.8, color="#A23B72")
axes[2].set_title(f'Test Set: MAPE = {best_rf_results["test_mape"]:.2f}%', fontsize=13, fontweight='bold')
axes[2].set_xlabel('Sample Index', fontsize=11)
axes[2].set_ylabel('Price (Rs)', fontsize=11)
axes[2].legend(fontsize=10)
axes[2].grid(alpha=0.3)

plt.suptitle(f'{best_rf_name}: Actual vs Predicted Carrot Prices',
             fontsize=16, fontweight='bold', y=0.995)
plt.tight_layout()
plt.savefig('/content/predictions_comparison.png', dpi=300, bbox_inches='tight')
plt.show()

print("✅ Predictions visualization saved")
```



Predictions visualization saved

```
# 3. Model Comparison Bar Charts
fig, axes = plt.subplots(1, 2, figsize=(16, 6))

# Extract numeric MAPE values
models = full_comparison['Model'].tolist()
mapes = [float(x.rstrip('%')) for x in full_comparison['Test MAPE'].tolist()]
r2s = [float(x) for x in full_comparison['Test R²'].tolist()]

# MAPE comparison
colors = ['forestgreen' if 'RF' in m or 'Gradient' in m else 'steelblue' for m in models]
axes[0].bar(range(len(models)), mapes, color=colors)
axes[0].set_xticks(range(len(models)))
axes[0].set_xticklabels(models, rotation=45, ha='right')
axes[0].set_ylabel('MAPE (%)', fontsize=12)
axes[0].set_title('Test MAPE Comparison\n(Lower is Better)', fontsize=14, fontweight='bold')
axes[0].grid(alpha=0.3, axis='y')
for i, v in enumerate(mapes):
    axes[0].text(i, v + 0.5, f'{v:.2f}%', ha='center', va='bottom', fontweight='bold')

# R² comparison
```

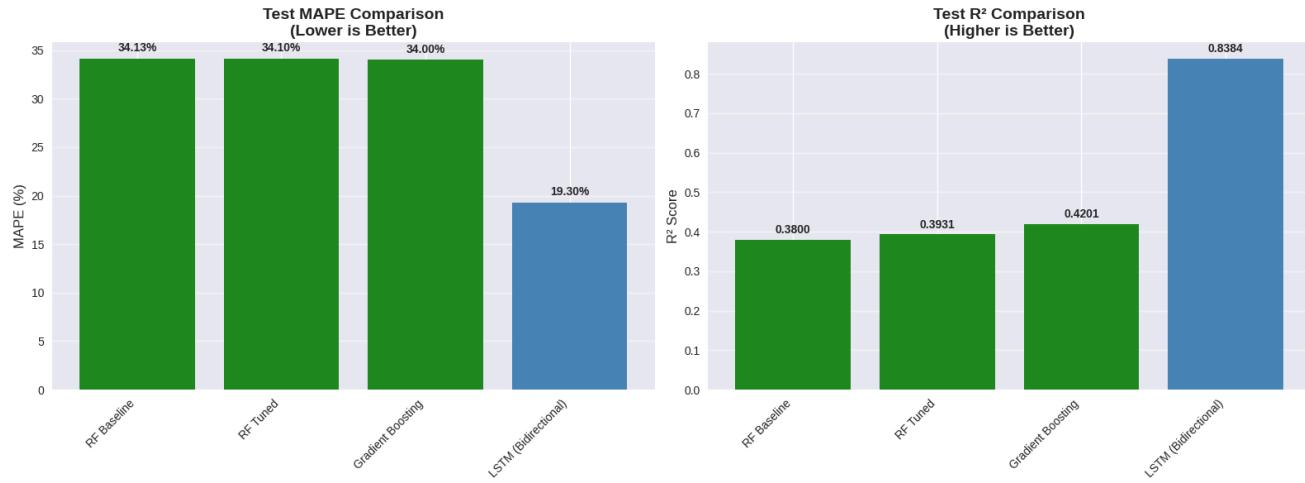
```

axes[1].bar(range(len(models)), r2s, color=colors)
axes[1].set_xticks(range(len(models)))
axes[1].set_xticklabels(models, rotation=45, ha='right')
axes[1].set_ylabel('R2 Score', fontsize=12)
axes[1].set_title('Test R2 Comparison\n(Higher is Better)', fontsize=14, fontweight='bold')
axes[1].grid(alpha=0.3, axis='y')
for i, v in enumerate(r2s):
    axes[1].text(i, v + 0.01, f'{v:.4f}', ha='center', va='bottom', fontweight='bold')

plt.tight_layout()
plt.savefig('/content/model_comparison_charts.png', dpi=300, bbox_inches='tight')
plt.show()

print("✅ Model comparison charts saved")

```



✅ Model comparison charts saved

## ▼ PART K: FINAL SUMMARY & CONCLUSIONS

```

print("=*80)
print("🎉 FINAL SUMMARY - RANDOM FOREST FOR CARROT PRICE PREDICTION")
print("=*80)

print("\n📊 DATASET:")
print(f" Total samples: {len(df_final)}")
print(f" Date range: {df_final.index.min()} to {df_final.index.max()}")
print(f" Features used: {len(final_features)}")
print(f" Train/Val/Test: {len(X_train)}/{len(X_val)}/{len(X_test)} samples")

print("\n🌲 MODELS TRAINED:")
print(" 1. Random Forest - Baseline")
print(" 2. Random Forest - Tuned (Hyperparameter Optimization)")
print(" 3. Gradient Boosting (Bonus)")

print("\n🎯 PERFORMANCE:")
print(full_comparison.to_string(index=False))

print(f"\n🏆 BEST MODEL: {best_rf_name}")
print(f" Test MAPE: {best_rf_mape:.2f}%")
print(f" Test R2: {comparison_df.loc[best_rf_idx, 'Test R2]}"})

print("\n📊 FEATURE IMPORTANCE (Top 5):")
for i, row in feature_imp_df.head(5).iterrows():
    print(f" {row['feature']}: {row['importance']:.4f}")

print("\n📈 COMPARISON TO LSTM:")
print(f" Best LSTM (Bidirectional): 19.30% MAPE, 0.8384 R2")
print(f" Best Random Forest: {best_rf_mape:.2f}% MAPE, {comparison_df.loc[best_rf_idx, 'Test R2]}) R2")

```